

# Algorithm Overview

The Boyer–Moore Majority Vote algorithm efficiently identifies a majority element in a sequence—an element that appears more than  $\lfloor n/2 \rfloor$  times. It works by maintaining a candidate and a counter while scanning the array once. Whenever the counter drops to zero, the current element becomes the new candidate. Finally, a verification pass ensures whether the candidate is indeed the majority. Its theoretical foundation lies in canceling out pairs of different elements, leaving only the potential majority candidate.

## Complexity Analysis

The Boyer–Moore algorithm runs in linear time  $O(n)$ , as it performs a single pass through the input array. The verification step adds another linear pass, resulting in an overall complexity of  $O(n)$ . Space complexity is  $O(1)$ , as it uses only a constant number of variables (candidate and counter). Best, average, and worst cases are all  $\Theta(n)$ , since each element is processed once. Compared with algorithms like hash-based counting ( $O(n)$  time,  $O(n)$  space), Boyer–Moore is optimal in space. In contrast, sorting-based approaches have  $O(n \log n)$  complexity, making Boyer–Moore more efficient for majority detection.

## Code Review

The code in `BoyerMooreMajorityVote.java` is well-structured and adheres to standard Java conventions. However, some inefficiencies exist: 1. Input validation could be improved to handle null or empty arrays early. 2. Repeated verification loops can be merged for minor optimization. 3. Memory footprint is minimal but could be improved by avoiding temporary list creation. Suggested optimizations include early termination when the majority count exceeds  $n/2$  and optional parallelism for large datasets. These could slightly improve practical runtime while maintaining  $O(n)$  asymptotic behavior.

## Empirical Results

Experimental runs show linear scaling of runtime with input size, confirming  $O(n)$  behavior. Performance plots (time vs. input size) demonstrate nearly constant per-element processing time. Theoretical predictions align closely with observed data, with constant factors primarily determined by JVM overhead and array traversal efficiency. Practical tests on arrays up to  $10^6$  elements show sub-second execution on modern hardware.

## Conclusion

The Boyer–Moore Majority Vote algorithm is optimal for majority detection due to its  $O(n)$  time and  $O(1)$  space complexity. The implementation is clean and effective, though minor optimizations can improve edge-case handling and practical speed. Empirical results validate theoretical expectations, making the algorithm an excellent balance between simplicity and efficiency.